

Ayush Suryawanshi

 suryawanshiayush007@gmail.com  +91-7000580565  [LinkedIn](#)

Skills

- **Languages:** Python, Go, JavaScript, SQL
- **Frameworks:** Django, Flask, FastAPI, Angular, LangGraph, MCP.
- **Databases:** PostgreSQL, MySQL, SQLite, MongoDB, Redis, ChromaDB, Milvus.
- **DevOps:** Git, Github, Docker, Kubernetes, Gunicorn, Nginx, Linux (Alma & Ubuntu), AWS (EC2, Sagemaker, Bedrock, RDS, Elasticsearch, S3, Lambda, VPC, IAM), MLFlow, Prometheus, Grafana.
- **Machine Learning:** Scikit Learn, OpenCV, XGBoost, YOLOv9.

Professional Experience

Senior Software Engineer | Innoviti Technologies – Bengaluru

September 2024 – Present

- Built and optimized a high-performance **inference stack** on **96 GB VRAM NVIDIA GPUs** with **CUDA 12.7** using **vLLM**, enabling **Qwen3-VL (32B, 8-bit quantized)** to serve **50+ concurrent requests at 3,000+ tokens/second** throughput with **sub-3-second latency**.
- Built AI Analyst, a context-aware **NLQ-to-SQL engine** handling **1,600+ queries/day** at **92% accuracy**, backend engineered with Python and **FastAPI**, deployed on **AWS EC2** with **Claude 3.7 Sonnet**, with integration to internal **PostgreSQL** servers for data querying.
- Engineered a Freshdesk ticket QC automation pipeline using **FastAPI** and **Celery** for **async processing**, integrating **Qwen3-VL 32B (8-bit quantized)** and **YOLOv9** to auto-validate support evidence across **5,000+ image artifacts daily** — reducing validation time from **1260 seconds to 7 seconds per ticket**.
- Engineered an automated **KYC verification** pipeline using Python and **Celery** for **async processing**, powered by **Qwen3-VL 32B (8-bit quantized)** — processing **900+ artifacts** daily at **96%+ accuracy** and cutting operational costs by **\$20,000 annually**.
- Designed a geo-optimized ASO dispatch algorithm in Python using **Haversine distance** and **priority-weighted allocation** against a **PostgreSQL** backend — enabling faster field routing and driving a **37% reduction in operational backlog**.
- Engineered the **backend** and **data pipeline** for Link — a fraud detection system processing **100,000+ transactions daily** across a **15,000+ merchant base**, built with **Python** and **XGBoost**, resulting in a **78% reduction in LEA complaints**.
- Architected Sheldon, a text-to-configuration automation engine powered by **Gemini 1.5 Flash** that dynamically generates Python code via LLM orchestration to **parse unstructured brand email** and Excel data into **PostgreSQL-compatible formats** — achieving **99.42% scheme accuracy** and **eliminating \$56,000 in annual misconfiguration losses**.

Software Development Engineer | Pheonix Solutions – Bengaluru

July 2022 – September 2024

- Developed **Disha-AI**, an address cleanup API using **FastAPI** — leveraging **Google Maps API** for **geocoding** and **Gemini 1.5 Flash** for intelligent **address normalization**, achieving **94%+ accuracy** across **20,000+ records** with **Redis caching** and **IP-level rate limiting**, **reducing validation costs by \$28,000**.
- Built a high-throughput inventory management system on **Django REST Framework + PostgreSQL** with **Redis caching layer** (write-through + TTL-based invalidation), sustaining **12K req/s peak load** across **50K+ stock units** with DB Optimization & connection pooling.
- Architected a production-grade compliance portal using **Django** and **Python** — integrating **OCR-driven license validation** and automated MVR compliance checks with **pHash** and **Celery async processing**, improving system throughput and **cutting review time by 30%**.
- Orchestrated a **Selenium-based** social profile crawler with **Django** backend and **Mistral 7B** for intelligent data extraction — processing **1,500+ profiles daily** at **92%+ accuracy** to enable optimal customer profiling.

Intern | KPMG India – Ahmedabad

Jan 2022 – June 2022

- Engineered an **Automated License Analyzer** for KPMG's Software Asset Managing Platform — parsing client licenses via **EasyOCR** and **PyPDF**, performing ETL in **python** and loading clean data into **SQLite**, **cutting analysis time by 90%**.
- Analyzed \$7 million worth of client license data using **python**, **pandas**, and **parquet** — delivering actionable insights through a live **Power BI dashboard**.

Gujarat Technological University, Ahmedabad

July 2018 – June 2022

- Bachelor of Engineering | 8.73 CGPA

Foundations & Continuous Growth

- Applied Study — **Designing Data-Intensive Applications (Martin Kleppmann)** - deep focus on distributed databases, scalable system design, fault tolerance, and consistency models. Concepts directly applied across production systems.
- Data Structures & Algorithms — Google course with strong focus on problem solving, complexity analysis, and algorithmic thinking.
- **LeetCode** — **200+ problems** solved across data structures, algorithms, and problem-solving patterns.